

Research Interests

As a systems researcher, I design algorithms that address problems at the intersection of networking, security and data science areas to build and deploy easy-to-use and scalable real-world systems.

Education

Summer 2018 **Princeton University**, *Ph.D.*, Computer Science.

(Expected) Advisor: Nick Feamster

How can networks run themselves? My dissertation research centers around answering this question. My goal is to design and build self-driving networks that satisfy the increased security, availability, and performance requirements of modern networks, by automatically (without humans in the loop) make holistic (not protocol-specific) control decisions in real time.

Spring 2013 **NC State University**, *M.S.*, Computer Science.

Spring 2009 **Indian Institute of Technology, Roorkee**, *B.Tech.*, Electronics & Comm.

Awards

2017 Best Paper Award winner, ACM SOSR

2017 Facebook Fellowship finalist

2016 Community Award winner, USENIX NSDI

2016 Juniper/Comcast SDN Throwdown winner

2015 Facebook Fellowship finalist

2013 Internet-2 Innovation Award winner

2013 Meissner Fellowship (Purdue University) winner

2010 College of Engineering Fellowship (NC State University) winner

Professional Experience

2015–Present **Princeton University**, *Research Assistant*, Princeton, NJ.

Mentors: Nick Feamster and Jennifer Rexford

Designed and implemented: a network streaming telemetry system, *Sonata* [10, 1, 7]; and an industrial-scale software-defined Internet exchange platform, *iSDX* [3, 2, 9, 8].

Summer 2016 **Microsoft Research**, *Research Intern*, Redmond, WA.

Mentors: Ratul Mahajan and Monia Ghobadi

Explored the design a wide-area network controller, *Roshan*, that configures both the optical (physical) and the network layer to make optimal use of limited available resources under failures.

2013–2014 **Georgia Tech**, *Research Assistant*, Princeton, NJ.

Mentor: Nick Feamster

Designed and built: a software-defined Internet exchange platform, *SDX* [5]; an event based network management tool, *Kinetic* [4]. Also, analysed multiple active and passive measurement dataset to model ISP interconnectivity in developing regions [12].

2011–2012 **NC State University**, *Research Assistant*, Raleigh, NC.

Mentor: Injong Rhee

Developed *WiFox* [6], solving the problem of performance degradation for large audience environments. This technology has been licensed out to Intel.

- Summer 2011 **Google, Research Intern**, Mountain View, CA.
Mentor: Nandita Dukkupati
Worked on quantifying the role played by TCP time outs on Google's search traffic. Instrumented the TCP stack for Google's front end servers to collect the data required for this measurement study.
- Spring 2010 **Indian Institute of Science, Project Assistant**, Bangalore, India.
Mentor: Anurag Kumar
Designed and implemented a WiFi AP based scheduling algorithm ensuring fairness to clients with disparate link qualities.

Publications

Conferences

- [1] **Arpit Gupta**, Rob Harrison, Ankita Pawar, Marco Canini, Nick Feamster, Jennifer Rexford, and Walter Willinger. Sonata: Query-Driven Network Telemetry. *Under Submission*.
- [2] Robert MacDavid, Rüdiger Birkner, Ori Rottenstreich, **Arpit Gupta**, Nick Feamster, and Jennifer Rexford. Concise Encoding of Flow Attributes in SDN Switches. In *ACM Symposium on SDN Research (SOSR)*, 2017.
Best Paper Award (1 out of 77).
- [3] **Arpit Gupta**, Robert MacDavid, Rüdiger Birkner, Marco Canini, Nick Feamster, Jennifer Rexford, and Laurent Vanbever. An Industrial-Scale Software Defined Internet Exchange Point. In *USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 2016.
Community Award (1 out of 255).
- [4] Hyojoon Kim, Joshua Reich, **Arpit Gupta**, Muhammad Shahbaz, Nick Feamster, and Russ Clark. Kinetic: Verifiable Dynamic Network Control. In *USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 2015.
60 citations till Nov 2017 based on Google Scholar.
- [5] **Arpit Gupta**, Laurent Vanbever, Muhammad Shahbaz, Sean Patrick Donovan, Brandon Schlinker, Nick Feamster, Jennifer Rexford, Scott Shenker, Russ Clark, and Ethan Katz-Bassett. SDX: A Software Defined Internet Exchange. In *ACM SIGCOMM*, 2014.
210 citations till Nov 2017 based on Google Scholar.
- [6] **Arpit Gupta**, Jeongki Min, and Injong Rhee. Wifox: Scaling wifi performance for large audience environments. In *ACM Conference on Emerging Networking Experiments and Technologies (CoNEXT)*, 2012.
50 citations till Nov 2017 based on Google Scholar.

Workshops & Short Papers

- [7] Rob Harrison, Cai Qizhe, **Arpit Gupta**, and Jennifer Rexford. Network-wide heavy hitter detection with commodity switches. *Under Submission*.
- [8] Xiaohe Hu, **Arpit Gupta**, Aurojit Panda, Nick Feamster, and Scott Shenker. Preserving Privacy at IXPs. *Under Submission*.
- [9] Rüdiger Birkner, **Arpit Gupta**, Nick Feamster, and Laurent Vanbever. SDX-Based Flexibility or Internet Correctness?: Pick Two! In *ACM Symposium on SDN Research (SOSR)*, 2017.
- [10] **Arpit Gupta**, Rüdiger Birkner, Marco Canini, Nick Feamster, Chris Mac-Stoker, and Walter Willinger. Network Monitoring as a Streaming Analytics Problem. In *ACM Workshop on Hot Topics in Networks (HotNets)*, 2016.

- [11] **Arpit Gupta**, Nick Feamster, and Laurent Vanbever. Authorizing Network Control at Software Defined Internet Exchange Points. In *ACM Symposium on SDN Research (SOSR)*, 2016.
- [12] **Arpit Gupta**, Matt Calder, Nick Feamster, Marshini Chetty, Enrico Calandro, and Ethan Katz-Bassett. Peering at the Internet's Frontier: A First Look at ISP Interconnectivity in Africa. In *Passive and Active Network Measurement (PAM)*, 2014.
55 citations till Nov 2017 based on Google Scholar.

Presentations

Sonata: Query-Driven Streaming Network Telemetry

Conferences: ACM HotNets (11/16), NANOG 70 (05/17), P4 Workshop (05/17)

Industry: Comcast (12/16), NIKSUN Inc. (06/17), AT&T (10/17)

Universities: Boston University (10/16)

iSDX: An Industrial-Scale Software Defined IXP

Conferences: USENIX NSDI (03/16), USENIX ATC (06/16), GENI Network Innovators Community Event (12/16)

Industry: AT&T (10/15), Project Endeavour (10/15), Corsa (11/15), CloudRouter (01/16), Open Networking Foundation Webinar (04/16), Appfest (05/16)

Universities: USC, Networked Systems Laboratory (08/15)

Authorizing Network Control at Software Defined IXPs

Conferences: ACM SOSR (03/16)

Industry: Verisign Inc. (08/15)

SDX: A Software Defined Internet Exchange

Conferences: ACM SIGCOMM (08/14), GENI Engineering Conference 20 (06/14), NANOG 59 (10/13), OpenIX Summit (04/15)

Industry: Facebook Inc. (08/14), Microsoft (08/14)

Universities: Stanford NetSeminar (10/14)

Peering at the Internet's Frontier

Conferences: Workshop on Passive and Active Measurements (03/14)

WiFox: Scaling WiFi Performance for Large Audience

Conferences: ACM SIGCOMM CoNEXT (12/12)

Universities: Duke University (10/12), UNC Chapel Hill (10/12)

Professional Activities

External Reviewer

NSDI 2014, ICNP 2016, SIGCOMM 2017, IEEE/ACM Transactions on Networking, IEEE Transactions on Mobile Computing, Computer Networks

Program Committee

Workshop on Self-Driving Networks, SIGCOMM 2018

Panelist

GENI Network Innovators Community Event 2016, CITP Conference on Global Internet Interconnection 2016

Teaching and Mentoring Experience

Teaching Assistant

- Spring 2016 Computer Networks (COS 461), Princeton University
Summer 2013 Computer Organization & Assembly Language, (CSC 236), NC State University
Fall 2012 Internet Protocols (CSC 573), NC State University

Course Development

- Spring 2016 Securing Cyberspace with Big Data (COS 598E), Princeton University
Summer 2015 Software Defined Networking, Coursera
Fall 2014 Software Defined Networking (CS 4270), Georgia Tech
Summer 2014 Software Defined Networking, Coursera

Guest Lecture

- Fall 2017 Computer Networks (COS 561), Princeton University
Spring 2017 Computer Networks (COS 561), Princeton University

Mentoring

- Fall 2017 – * David Liu, Ph.D.
Summer 2017 – * Bridger Hahn, M.S.
Fall 2016 – * Rob Harrison, Ph.D.
2015 – 2017 Rüdiger Birkner, M.S.
2015 – 2016 Robert MacDavid, M.S.

References

Prof. Jennifer Rexford
Department of Computer Science
Princeton University
35 Olden Street, CS 306
Princeton, NJ 08540
jrex@cs.princeton.edu

Prof. Nick Feamster
Department of Computer Science
Princeton University
310 Sherrerd Hall
Princeton, NJ 08540
feamster@cs.princeton.edu

Prof. Ethan Katz-Bassett
Department of Electrical Engineering
Columbia University
500 West 120th Street
New York, NY 10027
ethan@ee.columbia.edu

Dr. Walter Willinger
NIKSUN Inc.
457 N Harrison St
Princeton, NJ 08540
wwillinger@niksun.com

Research Statement

Arpit Gupta

As a **systems researcher**, I design algorithms that address problems at the intersection of **networking**, **security** and **data science** areas to build and deploy easy-to-use and scalable **real-world** systems. First, I discover real-world problems that network operators and vendors face while working “in the trenches” and understand their practical constraints. I then design innovative and provable solutions to these problems using various theoretical techniques (*e.g.*, optimization theory, and graph theory) and emerging technologies (*e.g.*, programmable data plane, scalable stream processors). Finally, I build realistic, easy-to-use, and scalable systems with minimal deployment overhead. My software packages are widely used in both academia and industry. For example, my software-defined Internet exchange point (IXP) prototype, which won the **Community Award** at *USENIX NSDI* in 2016, is used by many IXP operators across the globe. Also, the recently released, “network streaming telemetry” prototype, is currently being used by network operators and researchers at AT&T.

How can networks run themselves? My research centers around answering this question. My goal is to design and build self-driving networks that satisfy the increased security, availability, and performance requirements of modern networks, by automatically (without humans in the loop) make holistic (not protocol-specific) control decisions in real time. Although the research community explored the idea of self-driving networks in the past, there have been no notable attempts at building practically deployable solutions that satisfy current networks’ requirements.

Self-driving networks need to complete the control-loop where they need to collect, analyse, and react to various network events by themselves. My dissertation research focuses on the expressiveness, scalability, and correctness of two critical building blocks for this control loop. They are:

- *Telemetry* [2, 4, 7], which refers to the process of collecting and analyzing the raw network data to discover various network activities of interest in real time. For example, inferring presence of DDoS attacks and link failures in real time; and
- *Programmatic Control* [1, 3, 5, 6, 9, 10], which refers to the process of applying the fine-grained reactive actions without disrupting networks’ default behavior. For example, redirecting traffic from congested peering links without creating any forwarding loops.

Telemetry with *Sonata*

Problem. For networks to run themselves, they need to monitor a wide range of network activities. For example, they need to concurrently detect whether the network is under attack and also determine whether there is a device failure in the network. This involves extracting multiple features from the traffic data and combining them together to infer activities of interests in real time. Most existing real-time telemetry systems are either not expressive, *i.e.*, they support an insufficient set of telemetry tasks, or they are not scalable, *i.e.*, they fail to scale as the traffic volume intensifies and the number of telemetry tasks increases. In contrast, self-driving networks require running *multiple expressive* telemetry queries over *high-volume* networks. My research focuses on building a distributed real-time network telemetry system that is both *expressive* and *scalable*.

Observation. Many existing telemetry solutions take advantage of either scalable stream processing or programmable data-plane targets for network telemetry—but not both. While at first glance, these two technologies seem inherently different, they both apply a sequence of transformations over packets (tuples). This similarity provides new opportunities to combine their strengths. Since for most the telemetry tasks, the fraction of relevant traffic is tiny. Thus, unlike existing task (query)-agnostic systems, it makes sense to build query-driven telemetry systems. Based on these observations, I developed an expressive query-driven streaming telemetry system, *Sonata* [2, 4], that scales query execution by making use of both streaming and programmable data-plane targets and letting the output of queries themselves drive further processing.

Challenges and Solutions. Ideally, one would like to execute all queries in the data plane itself, but limited data-plane resources (*e.g.*, memory) restricts which portions of the queries to execute in the data plane. Requiring network operators to decide how to partition each query and configure the data-plane targets

to run the partitioned queries by themselves can be overwhelming. To address this problem, Sonata provides a declarative query interface where network operators express queries for a range of telemetry tasks as a sequence of dataflow operations over packet tuples. Internally, Sonata takes these queries, representative workload data, and various data-plane constraints to determine the partitioning plan for each query that minimizes the workload at the stream processor. It then compiles the partitioned queries to target-specific configurations. Sonata’s query planner also determines query-specific refinement plan that iteratively zooms-in over portions of traffic that satisfy the queries, making the best use of limited data-plane resources.

As a first step, I focused on a single-switch implementation for Sonata [2,4]; I’m currently extending this work to build a network-wide telemetry system. For the network-wide settings, Sonata’s query planner needs to determine: where to evaluate each query; how to partition the stateful operations over multiple hops while making the best use of limited resources along a path; and how to select the thresholds for filter operations [7].

Impact. Building an *expressive* and *scalable* telemetry system is a big step forward for self-driving networks. This project has generated a lot of interest in self-driving networks which is evident from two, recently announced, workshops (NSF and ACM SIGCOMM) on this topic. I have also open-sourced this project, and I am currently working with researchers at ISPs (AT&T) and security vendors (NIKSUN Inc.) to express and scale queries for various real-world telemetry tasks with Sonata.

Programmatic Control with SDX

Problem. After determining the reactive actions with telemetry, a self-driving network needs to apply these fine-grained actions in the data plane, usually in multi-domain settings. My research focuses on applying programmatic control for the inter-domain networks settings, where by default networks use border gateway protocol (BGP) to exchange traffic with each other. In my study, I demonstrate that BGP is not suited for applying fine-grained control actions. Ideally, one would like to replace BGP with a clean-slate solution supporting programmatic control at scale. However, the Internet-wide deployment of BGP-only routers makes such an approach impractical. Thus, the goal of my research was to design and implement a dirty-slate solution that ensures maximal impact with minimal deployment overhead while safely interoperating with BGP.

Observation. In recent years, we have seen the emergence of Internet exchange points (IXPs). They provide switching fabric for various networks for exchanging traffic with each other and are strategically located to influence a significant portion of the inter-domain traffic. I proposed building a software-defined exchange (SDX) [6] that replaces the conventional switching fabric with the programmable switches at IXPs. While simple in theory, my research shows that building a practically deployable SDX requires striking a delicate balance between expressiveness, correctness, privacy, and scalability.

Challenges and Solutions. Though SDX enables flexibility at the IXP, everyone else is still using BGP. Thus, it needs to ensure that this additional flexibility does not come at the cost of creating forwarding loops in the network. Requiring SDX participants to express their control program, while guaranteeing forwarding correctness, can be cumbersome, as they not only have to worry about their control program but also their peers’ programs. To address this problem, SDX provides a virtual switch abstraction to each participant which they can use to express their policies without worrying about others. SDX augments participant’s programs to ensure correctness. It then composes them together and compiles the composition to forwarding table entries for the data plane. While such an augmentation ensures correctness, the cost of additional forwarding rules in the data plane can create scalability problems. I developed novel “attribute encoding” schemes [5,6] that make use of several existing protocols and mechanisms to encode the forwarding attributes in the header of each packet before it enters the switching fabric—reducing the number of forwarding table entries required in the data plane.

To ensure that SDX participants can trust the computations at the IXP, I designed and implemented *SGDX* [8], that uses trusted execution environments such as Intel SGX to keep participants’ policies private. To ensure that SDX participants can only influence traffic that they are authorized to over shared physical switch, I built *FLANC* [3], which authorizes action requests over a shared switch for each participant. Finally, to stop participants at multiple IXPs from expressing conflicting policies, I developed *SIDR* [1], which makes it easier for participants to trade expressiveness (*i.e.*, fine-grained control policies) with privacy (*i.e.*,

share control policies with other SDXs) to ensure correctness.

Impact. My SDX paper [6] had a significant impact on the research community. It was the first project to demonstrate how programmable switches can enable flexible inter-domain routing at scale. It inspired NSF's first workshop on SDX in 2014, which was followed by **multi-million grants** for research in this area. The work also inspired the creation of *Endeavour*, a multi-university consortium in Europe for SDX-related research. To date, my SDX paper [6] has over 210 citations making it one of the **highest cited** papers from SIGCOMM 2014.

I open-sourced the *iSDX* [5] project with the Open Networking Foundation (ONF) and worked closely with two full-time software developers to make the SDX code production ready. I also worked with various network operators, switch vendors, and service providers to help deploy SDX at an inter-government agency exchange and many IXPs across the globe. Project *Endeavour* currently uses *iSDX* as the default platform for their projects. For these efforts, the project won the **Community Award** at *USENIX NSDI 2016*. Our *PathSets* [10] paper, that demonstrated how *iSDX*'s encoding scheme was applicable for more general network settings, won the **Best Paper** award at *SOSR 2017*.

Future Directions

Sonata and *SDX* are only the basic building blocks for self-driving networks. In future, I'd like to build-up on these works to continue working on the problems related to self-driving networks to make this distant dream a reality.

Expanding Networking Telemetry Footprint

Self-driving networks require insights from multiple vantage points in the network. So far, *Sonata* focuses primarily on inferring network activities at the core of the Internet (*e.g.*, border routers for large ISPs). Though such deployments provide broader visibility with limited overhead, they miss several insights such as end-to-end performance for specific web applications (*e.g.*, Netflix, CNN, etc.) in a region. These insights, available only at the edge of the Internet (*e.g.*, home routers, web browsers, etc.), are not only critical for robust self-driving networks but are also key to answering various policy questions on socio-political issues such as net neutrality, data privacy, censorship, etc.

Existing tools, such as BISmark, Ripe Atlas, etc., are not suited for edge telemetry as they are neither expressive or scalable. They either collect too much or too little data for analysis and are not suited for real-time query-driven telemetry. For example, one can use RIPE Atlas nodes to only express a limited set of telemetry queries tied to statically configured tools such as `ping`, `traceroute` etc. or use BISmark to collect packet traces only for limited duration and locations. I observed that it is possible to use cheaply available tools, such as programmable NICs (< \$40) and Raspberry Pis (< \$50) as data-plane and streaming targets respectively to build Sonata-like edge telemetry system.

Requiring network operators to express telemetry queries, while taking the computational capabilities of edge nodes into consideration and individually configure each edge node, can be overwhelming. Similar to Sonata's query interface, I envision providing a centralized query interface which abstracts away such details from the network operators. Under the hood, it determines a query plan that makes the best use of local compute resources before sending the packet tuples to a central processor.

To make this project more practical, I envision to design a wide-scale user study with *SonataEdge* to not only understand the real-world challenges for running it as scale, but also provides a platform for policy experts, journalists, and ISPs to answer various performance and policy related questions. For example, I am currently designing a wide-scale user study to quantify the relationship between subscribed uplink bandwidths and bit rates for streaming applications (*e.g.*, Netflix) for home networks.

Enabling Privacy-Preserving Network Telemetry

As we expand the footprint of network telemetry to build more robust self-driving networks, we encounter scenarios where consumers of the data are not the ones producing it. For example, end users generate quality-of-experience data for various latency-sensitive applications, but an eyeball ISP can consume this data to improve their network's performance. Though, in some cases, data consumers implicitly collect data in exchange for services they offer. However, currently, there are no solutions that let data producers explicitly share their data in a privacy-preserving manner. My goal is to enable data producers chose *what* portions of their data they'd like to share and with *whom* aiding networks in running themselves.

One approach is to share the data over multiple remote servers and rely on existing secure multi-party computation (SMPC) tools to ensure privacy-preserving computations. However, such an approach is query-agnostic and thus requires transferring all the (partitioned) raw data to remote servers for processing. An alternative is to process the data locally in a decentralized manner. I envision to build a system, which taps into the recent advances in the area of federated learning, where data consumers can execute their learning models over computing nodes that are local and private to end users. An end user can either use a dedicated device such as *SonataEdge*, or their mobile devices as local compute nodes. I also envision designing protocols that let remote data consumers train their models using these remote nodes to process the data locally in a privacy-preserving manner.

Developing such a decentralized system and protocols that are expressive, scalable, robust, and secure is non-trivial. My research agenda comprises of three inter-related thrusts that address these challenges: (1) *Expressive User Interface* for end users to express *what* data to trade and with *whom*. I envision collaborating with programming logic and privacy experts to create a domain-specific language for expressing data sharing policies; (2) *Data Privacy and Integrity* guarantees are required for data producers and consumers respectively. To guarantee data privacy, I plan to extend my previous work (*FLANC* [3]) to employ formal logic tools for generating policy compliance proofs for input telemetry queries. I also plan to design protocols that make use of TPMs at local nodes for guaranteeing data integrity; and (3) *Model Privacy* guarantees are required for data consumers. Learning models are the intellectual property of the consumers, and they might want to keep it private. To guarantee model privacy, I plan to design protocols that make use of existing works in the area of differential privacy and secure multi-party computation to ensure model privacy.

Co-designing Machine Learning and Query Planning Algorithms

Self-driving networks need to train learning models to discover activities of interests. Currently, learning algorithms only focus on training the learning models to maximize accuracy, and query planning algorithms solely focus on minimizing the execution cost for queries that represent these learning models. This disconnect between the two algorithms might result in very accurate, yet prohibitively expensive, learning models that cannot be deployed in the production networks.

I observed that it is possible to augment the optimization problem that most state-of-the-art learning algorithms solve to trade accuracy with the query execution cost while taking various resource constraints (e.g., limited memory in the data plane) to determine the learning model as well as the query plan. Based on this observation, I'd like to design and build telemetry systems, *SonataML*, that co-designs the algorithm for the machine learning and query planning problems to train realistic and deployable learning models.

Though the idea appears to be straightforward in principle, yet just embedding the cost and constraint metrics to the learning problems is not enough. The computational complexity of optimal solution grows exponentially with the increasing number of features for the learning problem. One needs to design heuristic algorithms that can find good enough solutions to this problem.

Another challenge is the availability of labeled training data. As most of the activities of interests are needles in a haystack, collecting labeled training data to bootstrap *SonataML* is extremely hard. I envision adapting existing active learning algorithms that can leverage the knowledge of the domain experts (security analysts and network operators) to generate as much of labeled data as possible while minimizing the workload for these human experts. As the nature of security and performance problems evolve with time, rather than replacing humans in one blow, I envision the active learning to be an intrinsic part of the self-driving networks, where the workload for the human experts reduces over time as the networks become more capable of running the networks by themselves.

To ensure that this solution is practically deployable, I plan to work with my collaborators at cloud (Google, Microsoft, etc.) and Internet (AT&T and Comcast) service providers, who have already shown great interest in this research direction. I also plan to work with my collaborators that make programmable switches (Barefoot Networks) to explore switch architectures that are better suited to support in-network machine learning.

Beyond Self-Driving Networks

The computational complexity of machine learning models has grown significantly in recent years. To catch up with this increased computational requirements, most cloud service providers have responded by

horizontally scaling the compute clusters. However, rather than mindlessly throwing more resources, one needs to step back and think more deeply to explore if there are new technologies on the horizon that can make better use of limited available resources.

I believe programmable data-plane targets provide the opportunity to vertically scale the analytics computation, *i.e.*, *Sonata's* idea of offloading query execution to data-plane targets applies to more general analytics settings. I plan to design and implement **general-purpose analytics stack** that is not only aware of the compute capabilities and limitations of CPU-based compute nodes but also the set of data-plane targets connecting these nodes.

To make this system a reality, I need to first demonstrate that using data-plane targets for offloading analytics computations is cheaper than using more CPU nodes. Answer what kind of data center topologies will make the best use of programmable data-plane targets? What kind of targets to use, *i.e.*, whether one should deploy cheaper but low throughput targets (*e.g.*, smart NICs), or high throughput but expensive targets (*e.g.*, switching fabrics), or both? Also, whether one should use off-the-shelf programmable data-plane targets or design custom switch architectures tailored to data analytics tasks? Second, I plan to design a query planner that for given topology, compute nodes and programmable data-plane targets; determines packet processing pipelines that execute over both the data-plane and the streaming targets. The goal of this query planner is to minimize the execution time while considering various data-plane constraints (*e.g.*, limited memory) and limitations (*e.g.*, cannot perform multiplication or division operation). The query planner also needs to determine the data structures for packet's payload—reducing the parsing overhead for the intermediate data-plane targets. Finally, I'd also like to expand the idea of co-designing the algorithms for learning and query planning by applying it to more commonly-used machine learning problems.

References

- [1] R. Birkner, A. Gupta, N. Feamster, and L. Vanbever. SDX-Based Flexibility or Internet Correctness?: Pick Two! In *ACM Symposium on SDN Research (SOSR)*, 2017.
- [2] A. Gupta, R. Birkner, M. Canini, N. Feamster, C. Mac-Stoker, and W. Willinger. Network Monitoring as a Streaming Analytics Problem. In *ACM Workshop on Hot Topics in Networks (HotNets)*, 2016.
- [3] A. Gupta, N. Feamster, and L. Vanbever. Authorizing Network Control at Software Defined Internet Exchange Points. In *ACM Symposium on SDN Research (SOSR)*, 2016.
- [4] A. Gupta, R. Harrison, A. Pawar, R. Birkner, M. Canini, N. Feamster, J. Rexford, and W. Willinger. Sonata: Query-Driven Network Telemetry. *arXiv preprint arXiv:1705.01049*, 2017.
- [5] A. Gupta, R. MacDavid, R. Birkner, M. Canini, N. Feamster, J. Rexford, and L. Vanbever. An Industrial-Scale Software Defined Internet Exchange Point. In *USENIX NSDI*, 2016.
- [6] A. Gupta, L. Vanbever, M. Shahbaz, S. P. Donovan, B. Schlinker, N. Feamster, J. Rexford, S. Shenker, R. Clark, and E. Katz-Bassett. SDX: A Software Defined Internet Exchange. In *ACM SIGCOMM*, 2014.
- [7] R. Harrison, C. Qizhe, A. Gupta, and J. Rexford. Network-Wide Heavy Hitter Detection with Commodity Switches. *Under Submission*.
- [8] X. Hu, **Arpit Gupta**, A. Panda, N. Feamster, and S. Shenker. Preserving Privacy at IXPs. *Under Submission*.
- [9] H. Kim, J. Reich, A. Gupta, M. Shahbaz, N. Feamster, and R. Clark. Kinetic: Verifiable Dynamic Network Control. In *USENIX NSDI*, 2015.
- [10] R. MacDavid, R. Birkner, O. Rottenstreich, A. Gupta, N. Feamster, and J. Rexford. Concise Encoding of Flow Attributes in SDN Switches. In *ACM Symposium on SDN Research (SOSR)*, 2017.

Teaching Statement

Arpit Gupta

As a graduate student, I have had the privilege to contribute to *eight* graduate, undergraduate, and on-line courses; and mentor *five* graduate students. I learned a lot from this experience. It strengthened my resolve to pursue a career in academia where I get the privilege not only to teach but also advise undergraduate and graduate students.

Lectures. I have given multiple lectures for the undergraduate and graduate networking classes at Princeton as a precept (COS 461 [2]) and guest lecturer (COS 461 [2] & COS 561 [3]). In my lectures, I pay special attention to instilling excitement about the material and help students discover new interests. For example, in my lectures, rather than speaking exclusively in the abstract, I motivated content-delivery networks with cat videos and used real-world configurations from Princeton’s network to explain firewalls and border routing. I believe a good lecture should discuss recent research advancements, showing students how the field is still evolving and how they can contribute to making an impact of their own.

Course Development. I have also contributed to the development of two new courses: **Software-Defined Networking** (SDN) and **Securing Cyberspace with Big Data** [4] at Georgia Tech (with Nick Feamster) and Princeton University (with Walter Willinger) respectively. The SDN course was offered both in-class [6] and online [5]. The online version, offered in Summer 2014 and 2015, drew more than *XX thousand* students.

Programming Assignment. Given the vast array of careers that students embark upon after graduation, it is crucial to provide them with transferable problem-solving skills. I believe programming assignments are the right medium to hone such skills. I have contributed to the design of multiple programming assignments for the networking courses at Princeton and Georgia Tech.

My approach is to translate cutting-edge research into programming assignments. It not only exposes students to latest developments in the area but also adds a new dimension to the research itself. For example, my experience designing programming assignments for the *SDX* [11] and the *Sonata* [12] projects forced me to make the programming interface for these systems more intuitive and easy-to-use. This experience was especially useful while training network operators—deploying my solutions in the production network.

Continuing Education. The Coursera course provided me with an opportunity to reach out to thousands of students from different age groups and with diverse backgrounds. Many of them were mid-career networking employees honing their skills to embrace the changing ecosystem, *i.e.*, shift to automated programmable networking. Being able to contribute to such a massive workforce retraining was a gratifying experience for me. This experience also helped my research. For example, we used one of the programming assignment as a user study quantifying how *Kinetic* [9] makes it easier for network operators to express stateful policies.

My teaching experiences confirmed that I enjoy teaching at all levels and over all mediums. I am looking forward to opportunities for teaching **computer networks**, **cybersecurity**, and **big-data systems** courses at the undergraduate and graduate levels. I am also interested in developing interdisciplinary graduate courses (with possibly *online* versions) that combine the areas of *networking*, *security*, and *data science*.

Advising. I was privileged to get an opportunity to mentor seven graduate students on specific projects. I supervised Robert MacDavid with Jennifer Rexford on the *iSDX* [8] and PathSets [10] projects. I mentored Rüdiger Birkner from ETH Zürich with Laurent Vanbever and Nick Feamster on the *iSDX* [8], *SIDR* [1], and *Sonata* [7] projects. I am currently mentoring Rob Harrison with Jennifer Rexford on the Sonata [7] project, Bridger Hahn with Nick Feamster on the IoT security project, and David Liu with Nick Feamster and Mike Freedman on co-designing machine learning and query planning algorithms. I have also mentored many junior students: Hooman Mohajeri, Disney Yan Lam, and Laura Roberts; at Princeton.

As a mentor, I try to help students discover their interests, help them identify problems where they can capitalize on their strengths, and develop a taste for research. For example, with Robert MacDavid (background in theory) and Rüdiger Birkner (background in systems), I worked on the design of the encoding algorithm and the design and implementation of production-quality prototype respectively for the *iSDX* project [8]. Giving them a well-defined problem, tailored to their strengths, at the beginning boosted their confidence to do research and make an impact. For example, Robert applied the techniques we developed for the *iSDX* project to more general settings (beyond Internet Exchange Points) and the project won the **best paper** award at ACM SOSR 2017 [10].

I will promote my students to interact with operators working “in the trenches” to learn about real-world problems. To facilitate this, I will encourage them to go out for industrial internships, especially during the early years of their graduate studies; and attend operational meet-ups and conferences, *e.g.*, North American network operator group (NANOG) meetings. Among the gamut of operational problems, I will provide constructive feedback to my students helping them to iteratively refine their ideas and identify impactful problems that require a more principled solution.

References

- [1] R. Birkner, A. Gupta, N. Feamster, and L. Vanbever. SDX-Based Flexibility or Internet Correctness?: Pick Two! In *ACM Symposium on SDN Research (SOSR)*, 2017.
- [2] COS 461: Computer Networks, Princeton University, Spring 2016. <https://www.cs.princeton.edu/courses/archive/spring16/cos461/>.
- [3] COS 561: Advanced Computer Networks, Princeton University, Fall 2017. <http://www.cs.princeton.edu/courses/archive/fall17/cos561/>.
- [4] COS 598: Securing Cyberspace with Big Data, Princeton University, Spring 2016. https://registrar.princeton.edu/course-offerings/course_details.xml?courseid=002132&term=1164.
- [5] Coursera SDN Course. class.coursera.org/sdn-002/.
- [6] CS 4270: Software Defined Networking, Georgia Tech, Fall 2014. <http://noise.gatech.edu/classes/cs8803sdn/fall2014/>.
- [7] A. Gupta, R. Harrison, A. Pawar, R. Birkner, M. Canini, N. Feamster, J. Rexford, and W. Willinger. Sonata: Query-Driven Network Telemetry. *arXiv preprint arXiv:1705.01049*, 2017.
- [8] A. Gupta, R. MacDavid, R. Birkner, M. Canini, N. Feamster, J. Rexford, and L. Vanbever. An Industrial-Scale Software Defined Internet Exchange Point. In *USENIX NSDI*, 2016.
- [9] H. Kim, J. Reich, A. Gupta, M. Shahbaz, N. Feamster, and R. Clark. Kinetic: Verifiable Dynamic Network Control. In *USENIX NSDI*, 2015.
- [10] R. MacDavid, R. Birkner, O. Rottenstreich, A. Gupta, N. Feamster, and J. Rexford. Concise Encoding of Flow Attributes in SDN Switches. In *ACM Symposium on SDN Research (SOSR)*, 2017.
- [11] SDX Programming Assignment. <https://github.com/PrincetonUniversity/Coursera-SDN/tree/master/assignments/sdn-ixp>.
- [12] Sonata Programming Assignment. <https://github.com/Sonata-Princeton/SONATA-DEV/tree/tutorial/sonata/tutorials/Tutorial-1>.